

legal-right-extractor - TDD - v1.2

Technical Design Document

Legal Rights Extraction System (Term Sheets / SHA)

Version 1.1 • Updated • Author: Corentin Marion

Goal

Build a **local-first, auditable extraction pipeline** that converts legal documents into a structured portfolio dataset of investor rights, each backed by explicit textual evidence and conflict flags.

The system prioritizes **determinism, traceability, and reproducibility** over coverage or abstraction.

Design Principles

- **Accuracy over coverage**
Prefer `present = false` over unsupported extraction.
 - **Auditability**
Every extracted item must include a direct quote and provenance (file, page, article when available).
 - **Deterministic scaffolding around the LLM**
Ingestion, chunking, retrieval, validation, conflict detection, and aggregation are explicit steps.
 - **Local-first by default**
Documents and models run locally; no mandatory external services.
-

Audience

- Engineers building applied extraction systems.
- Technically literate investment and legal operations users reviewing outputs.

1. System Overview

The system is a **pipeline controller** orchestrating deterministic stages and invoking an LLM **only for schema-locked extraction**.

The LLM does not decide, infer, or reconcile.
It extracts or flags uncertainty.

Architecture (MVP)

```
Ingestion
  ↓
  Chunking (page / article mapping)
  ↓
  Retrieval (BM25 lexical search)
  ↓
  Extraction (LLM, schema-locked)
  ↓
  Validation (deterministic rules)
  ↓
  Conflict Detection
  ↓
  Aggregation
  ↓
  Export / UI (CSV, XLSX, Streamlit)
```

Key Outputs

- Portfolio table (one row per `RightItem`)
- Per-company detail view with evidence excerpts (≤ 25 words)
- Run metadata for reproducibility

2. Open-Source Stack

All components are usable locally without paid APIs.

Layer	Tools	Rationale
LLM runtime	Ollama, llama.cpp, vLLM	Local inference, controllable

Layer	Tools	Rationale
Retrieval	BM25 (rank_bm25 / Whoosh / Lucene-style)	Exact lexical matching
PDF extraction	pdfplumber, pypdf	Page-level mapping
OCR (optional)	ocrmypdf + Tesseract	Fallback only
DOCX parsing	python-docx	Paragraph + heading extraction
Orchestration	Python + Pydantic	Schema enforcement
UI	Streamlit	Internal tooling
Export	pandas + openpyxl	Structured XLSX
Evaluation	promptfoo, Label Studio	Regression and gold sets

Embedding models and vector indexes are intentionally excluded in the current design.

3. Data Model and Schemas

All extracted rights are normalized into a single canonical object.

Core Object: RightItem

```

company_id: str
clause_family: enum[veto, liquidity, exit, anti_dilution]
clause_type: str
present: bool
beneficiary: enum[investor, founders, company, mixed, unclear]
trigger: str
effect: str
thresholds: str | null
evidence:
    quote: str      # ≤ 25 words
    source_file: str
    page: int | null
    article: str | null
confidence: float    # triage only
flags: list[enum[ambiguous, conflicting, missing_evidence]]
```

Normalization Rules

- One `RightItem` per distinct right.
 - No paraphrasing in evidence.
 - If evidence is missing, set `present = false` and flag `missing_evidence`.
-

4. Pipeline Design

4.1 Ingestion and Text Extraction

- Detect file type (PDF / DOCX / TXT).
- Store file hash for reproducibility.
- Extract per-page text for PDFs.
- Infer section/article structure where possible.
- OCR is a fallback, not default behavior.

Each extracted span retains `(source_file, page, char_start, char_end)`.

4.2 Chunking and Page Mapping

Chunking preserves legal structure.

- First split by detected articles/sections.
 - Then enforce a size cap (\approx 800–1200 tokens).
 - Maintain overlap (\approx 10–15%) only within sections.
 - Store metadata: `company_id`, `source_file`, `page_start`, `page_end`, `section_title`.
-

4.3 Retrieval (BM25)

Retrieval is **purely lexical**.

- Index all chunks using BM25.
- For each clause family, define a fixed query set (synonyms and drafting variants).
- Example (liquidity):
 - “right of first refusal”

- “transfer restriction”
- “lock-up”
- “preemption”

Top-k chunks (e.g. k = 8–15) are passed to extraction **with full provenance**.

No embeddings, no semantic similarity, no vector index.

5. Prompting and Extraction

Prompts are treated as versioned code.

The extractor:

- receives only retrieved chunks,
- outputs **valid JSON only**,
- follows the `RightItem` schema exactly.

Hard constraints:

- No inference.
 - No conflict resolution.
 - If unsure: `present = false` or `ambiguous`.
-

6. Validation and Aggregation

6.1 Deterministic Validation

- Reject `present = true` without evidence.
- Enforce quote length ≤ 25 words.
- Validate enums and required fields.
- Auto-flag low-confidence items.
- Fill missing page numbers deterministically if available upstream.

6.2 Conflict Detection

- Group by `(company_id, clause_family, clause_type)`.
- Compare normalized trigger / effect / thresholds strings.

- If differences exceed a strict lexical threshold, flag `conflicting`.
- Never override silently.

6.3 Aggregation and Export

- Aggregate validated `RightItem`s into a table.
 - Export CSV and XLSX.
 - Generate per-company drilldown views.
-

7. Security, Privacy, Reproducibility

- Confidential by default.
 - Minimal document retention.
 - No logging of full text.
 - Version prompts and schemas.
 - Store run metadata: model, prompt version, timestamp, git commit.
 - File hashing used to detect document changes.
-

8. Engineer Playbook

Step 0 — Environment

Install: pdfplumber, pypdf, python-docx, rank-bm25, pydantic, pandas, openpyxl, streamlit.

Step 1 — Ingestion

Extract per-page text and metadata.

Step 2 — Chunking

Apply rule-based legal segmentation.

Step 3 — Retrieval

Index chunks with BM25; query per clause family.

Step 4 — Extraction

Run schema-locked LLM prompts.

Step 5 — Validation + Conflicts

Apply deterministic rules.

Step 6 — Export / UI

Generate XLSX and Streamlit views.

Step 7 — Evaluation

Maintain gold sets and regression tests.

9. Design Rationale (Explicit)

BM25 was chosen over embeddings because:

- legal drafting relies on **exact language**,
- lexical transparency matters for audit,
- semantic similarity introduces uncontrolled recall,
- deterministic ranking is easier to reason about and debug.

This is a deliberate trade-off.